

# Model 1

- Optimiser = Adam
  - Epochs = 20
- 

## Model

```
[ ] x = inputs = keras.Input(shape=(32, 32, 3))

x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)

x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)

x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)

x = layers.Flatten()(x)
x = layers.Dense(64, activation="relu")(x)
x = layers.Dense(10)(x)

outputs = x
model = keras.Model(inputs=inputs, outputs=outputs)
```

- **Accuracy: 70%**
  - Model is from Kaggle
  - No Dropout / Augmentation
-

```
#####
# with dropout after each pooling layer
#####
# 3 dropout layers

x = inputs = keras.Input(shape=(32, 32, 3))

#1
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Dropout(0.25)(x)

#2
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Dropout(0.25)(x)

#3
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)

x = layers.Flatten()(x)
x = layers.Dense(64, activation="relu")(x)
x = layers.Dropout(0.25)(x)

x = layers.Dense(10)(x)

outputs = x
model = keras.Model(inputs=inputs, outputs=outputs)
```

- Same Model but with 3 drop out layers
  - One after each pooling layer
  - And one at the end before the last dense layer
  - 67% accuracy
  - Too many drop out layers
-

```
[8] #####
# dropout after each pooling layer
#####
# 2 dropout layers

x = inputs = keras.Input(shape=(32, 32, 3))

#1
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Dropout(0.25)(x)

#2
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Dropout(0.25)(x)

#3
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)

x = layers.Flatten()(x)
x = layers.Dense(64, activation="relu")(x)

x = layers.Dense(10)(x)

outputs = x
model = keras.Model(inputs=inputs, outputs=outputs)
```

- 2 dropout layers
  - 1 after each pooling layer
  - Removed the third dropout layer at the end
  - **Accuracy: 74%**
  - Removing the dropout at the end increases the accuracy
-



```
#####  
# dropout after each pooling layer  
#####  
# 2 dropout layers  
# Increased drop out to 0.5  
  
x = inputs = keras.Input(shape=(32, 32, 3))  
  
#1  
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)  
x = layers.MaxPooling2D(pool_size=2)(x)  
x = layers.Dropout(0.5)(x)  
  
#2  
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)  
x = layers.MaxPooling2D(pool_size=2)(x)  
x = layers.Dropout(0.5)(x)  
  
#3  
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)  
  
x = layers.Flatten()(x)  
x = layers.Dense(64, activation="relu")(x)  
  
x = layers.Dense(10)(x)  
  
outputs = x  
model = keras.Model(inputs=inputs, outputs=outputs)
```

- increase the 2 dropout layers to 0.5
  - **Accuracy 71%**
  - perhaps too much dropout
-

```
[12] #####
# dropout at the end
#####
# 1 dropout layers

x = inputs = keras.Input(shape=(32, 32, 3))

#1
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)

#2
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)

#3
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)

x = layers.Flatten()(x)
x = layers.Dense(64, activation="relu")(x)

x = layers.Dropout(0.25)(x)
x = layers.Dense(10)(x)

outputs = x
model = keras.Model(inputs=inputs, outputs=outputs)
```

- 1 dropout layer
  - at the end before the last dense layer
  - **Accuracy 72.6%**
-

```
[16] #####
# dropout at the end
#####
# 1 dropout layers
# increase the drop out ot 0.5

x = inputs = keras.Input(shape=(32, 32, 3))

#1
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)

#2
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)

#3
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)

x = layers.Flatten()(x)
x = layers.Dense(64, activation="relu")(x)

x = layers.Dropout(0.5)(x)
x = layers.Dense(10)(x)

outputs = x
model = keras.Model(inputs=inputs, outputs=outputs)
```

- 1 dropout layer
  - Increased the Dropout to 0.5
  - at the end before the last dense layer
  - Almost no increase in accuracy
  - **Accuracy 72.8%**
-

```
[32]
data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
    ]
)

x = inputs = keras.Input(shape=(32, 32, 3))
# 1
# Augmentation Layer
x = data_augmentation(inputs)
# 2
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
# 3
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
# 4
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)

x = layers.Flatten()(x)
x = layers.Dense(64, activation="relu")(x)
x = layers.Dense(10)(x)

outputs = x
model = keras.Model(inputs=inputs, outputs=outputs)
```

- No Dropout layers
  - 1 Augmentation layer
  - Random flip, Random rotation, random zoom
  - careful not to augment it too much
  - **Accuracy: 68%**
  - Possibly too much augmentation
-

```
[47]
data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomZoom(0.2),
    ]
)

x = inputs = keras.Input(shape=(32, 32, 3))
# 1
# Augmentation Layer
x = data_augmentation(inputs)
# 2
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
# 3
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
# 4
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)

x = layers.Flatten()(x)
x = layers.Dense(64, activation="relu")(x)
x = layers.Dense(10)(x)

outputs = x
model = keras.Model(inputs=inputs, outputs=outputs)
```

- Removed random rotation
  - **Accuracy 74%**
  - Nice improvement
  - If the image is augmented too much the useful features from the images are gone
-



```

data_augmentation = keras.Sequential(
    [
        layers.RandomTranslation(height_factor=(-0.1, 0.1), width_factor=(-0.1, 0.1)),
    ]
)

x = inputs = keras.Input(shape=(32, 32, 3))
# 1
# Augmentation Layer
x = data_augmentation(inputs)
# 2
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
# 3
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
# 4
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)

x = layers.Flatten()(x)
x = layers.Dense(64, activation="relu")(x)
x = layers.Dense(10)(x)

outputs = x
model = keras.Model(inputs=inputs, outputs=outputs)

```

- Only augmenting the width and height by 10%
- **Accuracy of 74%**

```

data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomZoom(0.2),
        layers.RandomTranslation(height_factor=(-0.1, 0.1), width_factor=(-0.1, 0.1)),
    ]
)


```

- Combine the two together
- **Accuracy 72%**
- Too much augmentation

```
data_augmentation = keras.Sequential(  
    [  
        layers.RandomTranslation(height_factor=(-0.1, 0.1), width_factor=(-0.1, 0.1)),  
    ]  
)  
  
x = inputs = keras.Input(shape=(32, 32, 3))  
# 1  
# Augmentation Layer  
x = data_augmentation(inputs)  
# 2  
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)  
x = layers.MaxPooling2D(pool_size=2)(x)  
x = layers.Dropout(0.25)(x)  
  
# 3  
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)  
x = layers.MaxPooling2D(pool_size=2)(x)  
x = layers.Dropout(0.25)(x)  
  
# 4  
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)  
  
x = layers.Flatten()(x)  
x = layers.Dense(64, activation="relu")(x)  
x = layers.Dense(10)(x)  
  
outputs = x  
model = keras.Model(inputs=inputs, outputs=outputs)
```

- Combine the best dropout model and the best augmentation model
- Model starts to underfit
- Accuracy of 73%

313/313 - 1s - loss: 0.7747 - accuracy: 0.7309 - 701ms/epoch - 2ms/step



```

data_augmentation = keras.Sequential(
    [
        layers.RandomTranslation(height_factor=(-0.1, 0.1), width_factor=(-0.1, 0.1)),
    ]
)

x = inputs = keras.Input(shape=(32, 32, 3))
# 1
# Augmentation Layer
x = data_augmentation(inputs)
# 2
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Dropout(0.25)(x)

# 3
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)

# 4
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)

x = layers.Flatten()(x)
x = layers.Dense(64, activation="relu")(x)
x = layers.Dense(10)(x)

outputs = x
model = keras.Model(inputs=inputs, outputs=outputs)

```

- Some Augmentation & 1 dropout layer
- Underfits
- **Accuracy: 67%**

---

## Summary So Far

- 2 Dropout layers 1 after each pooling layer got 74%
- Image augmentation can get also get up to 74%
- Augmenting the width/height, or the zoom and image flip produce the best result
- Have to be careful not to dropout too much or augment too much

# Takeaways

- Don't use too much dropout, careful where the dropout layers are. Too much dropout leads to underfitting

-